

Natural Language Processing Assignment

BODUGALA SATHVIK

2211CS020069

AIML-ALPHA

1. Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using edit distance and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

```
import re

import

pickle

import zlib

from collections import Counter


# Build corpus from a sample dictionary (you can enhance it with more words)

words = """going to china who was the first president of india winner of the match food in america"""


def words_list(text):

    return re.findall(r'\w+', text.lower())


WORDS = Counter(words_list( words)) #

Compression for large wordlist with

open('compressed_dict. pkl', 'wb') as f:
```

```

compressed = zlib.compress(pickle.dumps(WORDS))

f.write(compressed)

# Load dictionary in memory
def load_dictionary():

    with open('compressed_dict.pkl', 'rb') as f:

        return pickle.loads(zlib.decompress(f.read()))

def edit_distance_one(word):

    letters = 'abcdefghijklmnopqrstuvwxyz'

    splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]
    deletes = [L + R[1:] for L, R in splits if R]
    transposes = [L + R[1]
    + R[0] + R[2:] for L, R in splits if len(R) > 1]
    replaces = [L + c +
    R[1:] for L, R in splits if R for c in letters]
    inserts = [L + c + R for
    L, R in splits for c in letters]
    return set(deletes + transposes +
    replaces + inserts)

def known(words, dictionary):
    return set(w for w in words if w in dictionary)

def candidates(word, dictionary):

    return (known([word], dictionary) or
            known(edit_distance_one(word), dictionary) or
            [word])

def correct_word(word, dictionary):

    return max(candidates(word, dictionary), key=dictionary.get)

def correct_query(query, dictionary):

    return ''.join(correct_word(word, dictionary) for word in query.split())

# Main correction function

if __name__ == "__main__":

    dictionary = load_dictionary()
    n = int(input())
    queries = [input().strip() for _ in range(n)]

    for query in queries:
        print(correct_query(query, dictionary))

```

2. Deterministic Url and HashTag Segmentation

Explanation: This approach aims to find the most likely and meaningful segmentation of the input strings based on the provided dictionary of words and the constraint of selecting the longest valid tokens from the left.

Source Code:

```
import re

# Load words from words.txt into a set with
open("words.txt", "r") as file:

    dictionary = set(word.strip().lower() for word in file.readlines())


def is_number(s):

    """Check if the string is a number.""" try:

        float(s)

        return True

    except ValueError:

        return False


def tokenize(input_string, dictionary):

    """

    Tokenize the input string using the longest match first approach.

    Args:

        input_string: The string to be tokenized.

        dictionary: A set of valid words.

    Returns:

        A list of tokens from the input string.

    """

    length = len(input_string)

    if length == 0:

        return []
```

```

# dp[i] stores the tokens for the substring starting from index i
dp = [None] * (length + 1)
dp[0] = [] # Base case: empty string has no tokens

for i in range(1, length + 1):

    # Consider all possible ending positions for the current substring for
    j in range(i):

        left_part = input_string[j:i]

        # Check if left part is a valid word or number if
        (left_part in dictionary or is_number(left_part)) and
        ( dp[j] is not None
        ):

            # If left part is valid and remaining part has a valid
            tokenization      right_part_tokens      =      dp[j]

            right_part_tokens.append(left_part)

            # Choose the longest valid tokenization

            if len(right_part_tokens) > len(dp[i]) or dp[i] is None:

                dp[i] = right_part_tokens

# Return the tokenization for the entire string if it exists return
dp[length] if dp[length] is not None else [input_string]

def main():
    """Read input strings, tokenize them, and print the results.""" num_test_cases
    = int(input())

    for _ in range(num_test_cases):

        input_string = input().strip().lower()

        # Remove www and extensions for domain names, # for hashtags

        if input_string.startswith("www."):    input_string    =
        input_string[4:].rsplit(".", 1)[0] elif input_string.startswith("#"):

            input_string = input_string[1:]

```

```
tokens = tokenize(input_string, dictionary) print(f"Segmentation  
for Input: { ' '.join(tokens)}")
```

```
if __name__ == "__main__":
```

```
    main()
```

3. Disambiguation: Mouse vs Mouse

Explanation: This code provides a basic framework for classifying the usage of the word "mouse" in a sentence. You can further improve the accuracy by:

- a. **Expanding the Training Data:** Use a larger and more diverse dataset of sentences.
- b. **Experimenting with Different Classifiers:** Try other machine learning models like Support Vector Machines (SVM) or Random Forests.
- c. **Using Word Embeddings:** Consider using word embeddings like Word2Vec or GloVe to capture semantic relationships between words.

Source Code:

```
import pickle

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB

# Training data (sample corpus) training_sentences
= [
    "The complete mouse reference genome was sequenced in 2002.",
    "Tail length varies according to the environmental temperature of the mouse "
    "during postnatal development.",
    "A mouse is an input device.",
    "Many mice have a pink tail.",
    "The mouse pointer on the screen helps in navigation.",
    "A rodent like a mouse has sharp teeth.",
    "The mouse was connected to the computer using a USB port.",
    "The house was infested with mice.",
    "Computer users often prefer a wireless mouse."
]

# Labels corresponding to the training sentences
labels = [
    "animal",
    "animal",
```

```

"computer-mouse",
"animal",
"computer-mouse",
"animal",
"computer-mouse",
"animal",
"computer-mouse"
]

```

```

# Vectorize the training sentences vectorizer
= CountVectorizer()
X_train = vectorizer.fit_transform(training_sentences)

```

```

# Create and train the Naive Bayes classifier
classifier = MultinomialNB()
classifier.fit(X_train, labels)

```

```

# Function to predict the type of mouse def

```

```

predict_mouse_type(sentence):

```

```

"""

```

Predicts whether the 'mouse' in the sentence refers to an animal or a computer mouse.

Args:

 sentence: The input sentence.

Returns:

 "animal" or "computer-mouse"

```

"""

```

```

vectorized_sentence =
vectorizer.transform([sentence])    prediction =
classifier.predict(vectorized_sentence)[0]    return
prediction

```

```
# Get number of test cases num_test_cases
= int(input())
```

```
# Process each test case
for _ in range(num_test_cases):
    sentence = input() prediction =
    predict_mouse_type(sentence)
    print(prediction)
```

```
# Optionally, save the trained model for later use
with open('mouse_classifier. pkl', 'wb') as f:
    pickle.dump((vectorizer, classifier), f)
```

4. Language Detection

- Explanation:** This function loads the pre-trained model from a serialized file.
- It takes a text snippet as input, normalizes it to ASCII, and converts it into a TF- IDF vector using the loaded vectorizer.
- The function then uses the trained classifier to predict the language of the snippet based on the extracted features.

Source Code:

```
import pickle import unicodedata from
sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB

def normalize_to_ascii(text):
    """Remove non-ASCII characters and normalize text."""
    return unicodedata.normalize("NFKD", text).encode("ascii", "ignore").decode("ascii")

# Step 1: Training Data training_texts
= {
    "English": [
```



```

    "The quick brown fox jumps over the lazy dog.",
    "Rip Van Winkle is a story set in the years before the American Revolutionary War.",
],
"French": [
    "Le renard brun rapide saute par-dessus le chien paresseux.",
    "La revolution francaise a marque une periode importante de l'histoire.",
],
"German": [
    "Der schnelle braune Fuchs springt uber den faulen Hund.",
    "Die deutsche Wiedervereinigung war ein historisches Ereignis.",
],
"Spanish": [
    "El rapido zorro marron salta sobre el perro perezoso.",
    "La Revolucion Espanola fue un momento clave en la historia."
    "Si quieres que te asciendan te tienes que poner las pilas.",
],
}

```

```

# Normalize training data to ASCII labels

```

```

= []
texts = []
for language, samples in training_texts.items():
    labels.extend([language] * len(samples))
    texts.extend([normalize_to_ascii(sample) for sample in
samples])

```

```

# Step 2: Preprocessing and Feature Extraction vectorizer =

```

```

TfidfVectorizer(ngram_range=(2, 4), analyzer="char")

```

```

X_train = vectorizer.fit_transform(texts)

```

```

# Step 3: Train the Model

```

```

classifier =

```

```
MultinomialNB()
```

```
classifier.fit(X_train,
```

```
labels)
```

```
# Step 4: Serialize the Model
```

```
with open("language_model. pkl", "wb") as model_file:
```

```
    pickle.dump((vectorizer, classifier), model_file)
```

```
# Step 5: Language Detection Function def
```

```
def detect_language(snippet):
```

```
    with open("language_model. pkl", "rb") as model_file:
```

```
        vectorizer, classifier = pickle.load(model_file)
```

```
    # Normalize snippet to ASCII
```

```
    snippet = normalize_to_ascii(snippet)
```

```
    X_test =
```

```
    vectorizer.transform([snippet])
```

```
    prediction = classifier.predict(X_test)
```

```
    return prediction[0]
```

```
# Input Processing
```

```
if __name__ == "__main__": #
```

```
    Read multi-line input
```

```
    snippet = "" while
```

```
        True:
```

```
            try:
```

```
line = input() if
```

```
    line.strip(): snippet +=
```

```
    line + " " except
```

```
EOFError:
```

```
    break
```

```
# Predict and Output
```

```
detected_language = detect_language(snippet.strip())  
print(detected_language)
```

5. The Missing Apostrophes

Explanation Apostrophe Handling: The code defines a function `restore_apostrophes` that iterates through each word in the input text. It uses a combination of explicit checks for common contractions (e.g., "don't," "can't," "I've") and a regular expression to handle possessive nouns (e.g., "cat's," "dog's") to restore apostrophes where appropriate.

Source Code:

```
import re  
  
# Function to handle apostrophes for contractions and possessives def  
restore_apostrophes( text):  
    restored_text = [] words  
    = text.split()  
  
    for word in words:  
        lower_word = word.lower()  
  
        # Handle contractions if  
        word.lower() == "dont":  
            restored_text.append("don't")  
        elif word.lower() == "wont":  
            restored_text.append("won't")  
        elif word.lower() == "cant":  
            restored_text.append("can't")  
        elif word.lower() == "isnt":  
            restored_text.append("isn't")  
        elif word.lower() == "arent":  
            restored_text.append("aren't")  
        elif word.lower() == "wasnt":  
            restored_text.append("wasn't") elif  
        word.lower() == "werent":
```

```

        restored_text.append("weren't")
elif word.lower() == "hasnt":
    restored_text.append("hasn't")
elif word.lower() == "havent":
    restored_text.append("haven't")
elif word.lower() == "hadnt":
    restored_text.append("hadn't") elif
word.lower() == "didnt":
    restored_text.append("didn't")
elif word.lower() == "ive":
    restored_text.append("I've")
elif word.lower() == "were":
    restored_text.append("we're")
elif word.lower() == "i":
    restored_text.append("I")
elif word.lower() == "id":
    restored_text.append("I'd")
elif word.lower() == "ive":
    restored_text.append("I've")
elif word.lower() == "youve":
    restored_text.append("you've")
elif word.lower() == "hes":
    restored_text.append("he's")
elif word.lower() == "shes":
    restored_text.append("she's")
elif word.lower() == "its":
    restored_text.append("it's") elif
word.lower() == "were":
    restored_text.append("we're")

# Handle possessives (only add 's when it makes sense)
elif re.match(r'\w+s$', word) and lower_word not in ["its", "hers", "ours", "yours", "theirs"]:
    restored_text.append(re.sub(r"s$", "", word))

```

```

# For normal words that don't need apostrophes, keep them as is else:
    restored_text.append(word)

return " ".join(restored_text)

# Input

input_text = ""At a news conference Thursday at the Russian manned-space facility in Baikonur, Kazakhstan, Kornienko said "we will be missing nature, we will be missing landscapes, woods." He admitted that on his previous trip into space in 2010 "I even asked our psychological support folks to send me a calendar with photographs of nature, of rivers, of woods, of lakes."

Kelly was asked if hed miss his twin brother Mark, who also was an astronaut.

"Were used to this kind of thing," he said. "Ive gone longer without seeing him and it was great."

The mission wont be the longest time that a human has spent in space - four Russians spent a year or more aboard the Soviet-built Mir space station in the 1990s.

SCI Astronaut Twins

Scott Kelly (left) was asked Thursday if hed miss his twin brother, Mark, who also was an astronaut. Were used to this kind of thing, he said. Ive gone longer without seeing him and it was great. (NASA/Associated Press)

"The last time we had such a long duration flight was almost 20 years and of course al{-truncated-}"

# Restore apostrophes
output_text =
restore_apostrophes(input_text)
print(output_text)

```

6. Segment the Twitter Hashtags

Explanation: Tokenization with Dynamic Programming: The `segment_hashtag` function uses dynamic programming to break down the hashtag into a sequence of words. It iterates through the hashtag, checking for valid word combinations from a given dictionary and selecting the longest possible valid sequence.

Source Code:

```

# Define a function that segments a single hashtag into words
def segment_hashtag(hashtag, word_dict):
    n = len(hashtag)
    dp = [None] * (n + 1)

```

```

dp[0] = [] # Base case: empty string can be segmented as an empty list

# Iterate over the hashtag string
for i in range(1, n + 1):
    for j in range(max(0, i - 20), i): # Limit the length of words checked
        word = hashtag[j:i]
        if word in word_dict and dp[j] is not None:
            dp[i] = dp[j] + [word]
            break

return " ".join(dp[n]) if dp[n] is not None else hashtag

# Main function to process input and output results
def process_hashtags( num_hashtags, hashtags, word_dict):
    result = []
    for hashtag in hashtags:
        segmented = segment_hashtag(hashtag,
        word_dict) result.append(segmented) return result
# Sample dictionary of common words (expand this as needed) word_dict
= {
    "we", "are", "the", "people", "mention", "your", "faves",
    "now", "playing", "walking", "dead", "follow", "me"
}

# Sample input num_hashtags
= int(input())
hashtags = [input().strip() for _ in range(num_hashtags)]

# Process the hashtags and print the result
segmented_hashtags = process_hashtags(num_hashtags, hashtags, word_dict) for
segmented in segmented_hashtags:

```

```
print(segmented)
```

7. Expand the Acronyms

Explanation: Acronym Extraction: The code extracts acronyms and their potential expansions from a given set of text snippets by identifying uppercase words within parentheses and searching for preceding phrases. It also attempts to extract acronyms not explicitly defined in parentheses by analyzing the surrounding context.

Source Code:

```
import re

def extract_acronyms_and_expansions(snippets):
    """
    Extract acronyms and their expansions from the provided snippets.
    """
    acronym_dict = {}
    for snippet in snippets:
        # Find all potential acronyms (uppercase words typically enclosed in parentheses) matches
        matches = re.findall(r'\((\b[A-Z]+\b)\)', snippet)

        for match in matches:
            # Extract the preceding text (potential expansion) preceding_text
            preceding_text = snippet.split(f"({match})")[0].strip()

            # Look for the last meaningful phrase before the acronym
            expansion_candidates = re.split(r'[.,;:-]', preceding_text)
            if expansion_candidates:
                expansion = expansion_candidates[-1].strip()
                acronym_dict[match] = expansion

    # Additionally, handle acronyms not in parentheses but defined explicitly
    words = snippet.split()
    for i, word in enumerate(words):
        if word.isupper() and len(word) > 1: # Likely an acronym
            if word not in acronym_dict:
```



```

        # Try to extract its expansion from the surrounding context
        if i > 0:
            preceding_context = " ".join(words[max(0, i-5):i]) if preceding_context:
                acronym_dict[word] = preceding_context
        return acronym_dict

def process_tests(acronym_dict, tests):
    """
    Process test acronyms and return their expansions.
    """
    results = []
    for test in tests:
        # Normalize the test acronym (case insensitive)
        expansion = acronym_dict.get(test.upper(), "Not Found")
        results.append(expansion)
    return results

def main():
    # Read input
    n = int(input().strip())
    snippets = [input().strip() for _ in range(n)]
    tests = [input().strip() for _ in range(n)]

    # Extract acronyms and expansions
    acronym_dict = extract_acronyms_and_expansions(snippets)

    # Process test queries
    results = process_tests(acronym_dict, tests)

    # Output results
    print("\n".join(results))

if __name__ == "__main__": main()

```

8. Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using edit distance and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

```
import re

import

pickle

import zlib

from collections import Counter


# Build corpus from a sample dictionary (you can enhance it with more words)

words = """going to china who was the first president of india winner of the match food in america"""


def words_list(text):

    return re.findall(r'\w+', text.lower())


WORDS = Counter(words_list( words))


# Compression for large wordlist

with open('compressed_dict. pkl', 'wb') as f:

    compressed = zlib.compress(pickle.dumps(WORDS))

    f.write(compressed)


# Load dictionary in memory def

load_dictionary():

    with open('compressed_dict. pkl', 'rb') as f:

        return pickle.loads(zlib.decompress(f.read()))


def edit_distance_one(word):

    letters = 'abcdefghijklmnopqrstuvwxyz'
```

```

splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]
deletes = [L + R[1:] for L, R in splits if R]

transposes = [L + R[1] + R[0] + R[2:] for L, R in splits if len(R) >
1] replaces = [L + c + R[1:] for L, R in splits if R for c in letters]
inserts = [L + c + R for L, R in splits for c in letters]

return set(deletes + transposes + replaces + inserts)

def known(words, dictionary):
    return set(w for w in words if w in dictionary)

def candidates(word, dictionary):
    return (known([word], dictionary) or
            known(edit_distance_one(word), dictionary) or
            [word])

def correct_word(word, dictionary):
    return max(candidates(word, dictionary), key=dictionary.get)

def correct_query(query, dictionary):
    return ''.join(correct_word(word, dictionary) for word in query.split())

# Main correction function if
name__ == "__main__":
    dictionary = load_dictionary()
    n = int(input())
    queries = [input().strip() for _ in range(n)]

    for query in queries:
        print(correct_query(query, dictionary))

```

9. A Text-Processing Warmup

Explanation: Article and Date Counting: The code defines a function `count_articles_and_dates` that takes a text fragment as input. It first normalizes the text to lowercase for case-insensitive article counting. Then, it uses regular

expressions to count occurrences of the definite and indefinite articles ("a," "an," "the") and identify valid dates in various formats (e.g., "DD Month YYYY," "Month DD, YYYY," etc.)

Source Code:

```
import re

def count_articles_and_dates(fragment):
    """
    Count occurrences of 'a', 'an', 'the', and valid dates in a given text fragment.
    """
    # Normalize text for article counting
    lower_fragment = fragment.lower()

    # Count articles
    a_count = len(re.findall(r'\b[a]\b', lower_fragment))
    an_count = len(re.findall(r'\b[an]\b', lower_fragment))
    the_count = len(re.findall(r'\b[the]\b', lower_fragment))

    # Identify valid dates
    date_patterns = [
        r'\b\d{1,2}?(?:st|nd|rd|th)?(?:\s+of)?\s+(January|February|March|April|May|June|July|August|September|October|November|December)\s+\d{2,4}\b', # Day Month Year
        r'\b(January|February|March|April|May|June|July|August|September|October|November|December)\s+\d{1,2}?(?:st|nd|rd)?\s+\d{2,4}\b', # Month Day Year
        r'\b\d{1,2}/\d{1,2}/\d{2,4}\b', # Day/Month/Year
        r'\b\d{4}-\d{2}-\d{2}\b' # ISO format: Year-Month-Day
    ]
    # Combine all date patterns
    date_regex = '|'.join(date_patterns)
    dates = re.findall(date_regex, fragment, re.IGNORECASE)
    date_count = len(dates)
```

```

    return a_count, an_count, the_count, date_count

def main():
    import sys
    input = sys.stdin.read

    # Read input data
    data = input().strip().split("\n")
    t = int(data[0]) # Number of test cases
    fragments = data[1:] # Remaining lines contain the fragments

    results = []
    for i in range(t):
        fragment = fragments[i].strip() # Count articles and dates a_count, an_count,
        the_count, date_count = count_articles_and_dates(fragment)
        results.append(f"{a_count}\n{an_count}\n{the_count}\n{ date_count}")

    # Output results
    print("\n".join(results))

if __name__ == "__main__":
    main()

```

10. Who is it?

Explanation: Pronoun Identification and Entity Matching: The code first finds all pronouns (words enclosed in double backslashes) and their positions in the text. It then cleans the text by removing the backslashes. Next, it iterates through each pronoun and searches for the closest matching entity (from a provided list) that appears before the pronoun in the text.

Source Code:

```
import re
```

```

def resolve_pronouns(text, entities): #
    Extract all pronouns and their positions
    pronoun_pattern = r'\\(\\w+)\\'
    pronouns = [(match.group(1), match.start()) for match in re.finditer(pronoun_pattern, text)]

    # Clean the text by removing ** markers
    clean_text = re.sub(r'\\(\\w+)\\', r'1', text)

    # Initialize a list to store the resolved entities
    resolved = []

    # For each pronoun, find the corresponding
    entity for pronoun, pos in pronouns:
        closest_entity = None
        closest_distance = float('inf')

        # Iterate through all entities to find the best match
        for entity in entities:
            entity_pos = clean_text.rfind(entity, 0, pos) # Find the last occurrence of the entity before the pronoun

            if entity_pos != -1:
                distance = pos - (entity_pos + len(entity))
                if distance < closest_distance:
                    closest_distance = distance
                    closest_entity = entity

        # Append the resolved entity to the list
        resolved.append(closest_entity)

    return resolved

```

```
def main():

    import sys

    input = sys.stdin.read

    data = input().strip().split("\n")

    # Read the number of lines in the text snippet

    n = int(data[0])

    # Combine the next N lines into the full text snippet

    text_snippet = " ".join(data[1:n + 1])

    # Read the list of entities
    entities = [e.strip() for e in data[n +
1].split(';')]

    # Resolve pronouns

    result = resolve_pronouns(text_snippet, entities)

    # Output the resolved

    entities for entity in result:

        print(entity)

if __name__ == "__main__":

    main()
```